

---

### 3. MEASURING THE EFFECT OF A SUGAR TAX

## WORKING IN PYTHON

---

#### PYTHON-SPECIFIC LEARNING OBJECTIVES

In addition to the learning objectives for this project, in this section you will learn how to use method chaining to run a sequence of functions.

#### GETTING STARTED IN PYTHON

Head to the ‘Getting Started in Python’ (<https://tinyco.re/7775647>) page for help and advice on setting up a Python session to work with. Remember, you can run any page from this project as a *notebook* by downloading the relevant file from this repository and running it on your own computer. Alternatively, you can run pages online in your browser over at Binder.

#### *Preliminary settings*

Let’s import the packages we’ll need and also configure the settings we want:

```
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np
from pathlib import Path
import pingouin as pg
from lets_plot import *

LetsPlot.setup_html(no_js=True)
```

#### DOWNLOAD THE CODE

You can download the code chunks used in this project from the online version of this project (<https://tinyco.re/5461194>).

Don’t forget to also download the data into your working directory by following the steps in this project.

## PART 3.1 BEFORE-AND-AFTER COMPARISONS OF RETAIL PRICES

We will first look at price data from the treatment group (stores in Berkeley) to see what happened to the price of sugary and non-sugary beverages after the tax.

- Download the ‘Berkeley Store Price Survey’ Excel dataset, which is taken from the Global Food Research Program’s website.
- The first tab of the Excel file contains the data dictionary. Make sure you read the data description column carefully, and check that each variable is in the Data tab.

### PYTHON WALK-THROUGH 3.1

#### *Importing the data file into Python*

The data is stored in `.xlsx` format, so we use the `pd.read_excel` function from the `pandas` package to import it. We also import the `matplotlib` library as this includes packages that we will use later for drawing charts.

If you open the data in Microsoft’s Excel, Apple’s Numbers, or LibreOffice’s Calc (free), you will see that there are two tabs: ‘Data Dictionary’, which contains some information about the variables, and ‘Data’, which contains the actual data. Let’s import both into Python, so that we don’t need to refer back to Excel for the additional information about variables.

Remember to store the downloaded file in a subdirectory of your current working directory called `data`. Let’s import the data.

```
path_to_data = Path("data/Dataset Project 3.xlsx")

var_info = pd.read_excel(path_to_data, sheet_name="Data
Dictionary")
data = pd.read_excel(path_to_data, sheet_name="Data")
```

Let’s use the `info()` method to ensure that all the variables were given the correct datatypes:

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2175 entries, 0 to 2174
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   store_id         2175 non-null    int64  
 1   type             2175 non-null    object  
 2   store_type       2175 non-null    int64  
 ...
```

```

3   type2          235 non-null  object
4   size           2175 non-null float64
5   price          2175 non-null float64
6   price_per_oz   2175 non-null float64
7   price_per_oz_c 2175 non-null float64
8   taxed          2175 non-null int64
9   supp           2175 non-null int64
10  time           2175 non-null object
11  product_id    2175 non-null int64
dtypes: float64(4), int64(5), object(3)
memory usage: 204.0+ KB

```

Python classified all the variables containing numbers as either `int64` (integers) or `float64` (real numbers). However, other variables were labelled as `object` because they don't contain numbers, and some of the columns labelled as numbers are actually categorical variables (there are a fixed number of categories).

In the next section, we're going to use a 'for' loop. This loop has the syntax: `for thing in things`, then a colon (`:`), then an indented expression that does some operation on `thing`.

We are using a 'for' loop to redefine a list of incorrectly defined variables (such as `taxed`) as categorical data types. Rather than laboriously run the same code multiple times (`data["type"] = data["type"].astype("category")` and then `data["taxed"] = data["taxed"].astype("category")`, and so on), we can just have a list of variables we'd like to convert and then use a 'for' loop to apply the same kind of command to each variable. In the code below, `col` is our 'thing' (according to the 'for' loop syntax). When Python gets to `data[col] = data[col].astype("category")` it will replace `col` with each element in the list of things we feed it in turn.

So now let's ensure that `type`, `taxed`, `supp`, `store_id`, `store_type`, `type2` and `product_id` are represented as categorical variables. We'll use the `astype` function to convert them within a 'for' loop.

```

for col in ["type", "taxed", "supp", "store_id",
"store_type", "type2", "product_id"]:
    data[col] = data[col].astype("category")

# rename some of the categories
data["taxed"] = data["taxed"].cat.rename_categories(["not
taxed", "taxed"])
data["supp"] =
data["supp"].cat.rename_categories(["standard",
"supplemental"])
data["store_type"] =
data["store_type"].cat.rename_categories(

```

```

        ["Large supermarket", "Small supermarket", "Pharmacy",
"Gas station"]
)

```

You can see that we used `.cat.rename_categories` to give more meaningful names to different categories (where they are clearly defined).

There is another variable, `time`, which is classified as having an `object` datatype, but should be a categorical variable. Before we do this, we use the `unique` command to check the categories of this variable.

```
data["time"].unique()
```

```
array(['DEC2014', 'JUN2015', 'MAR2015'], dtype=object)
```

If you look at the timeline in the Silver *et al.* (2017) paper, you will notice that the third price survey was in March 2016, not in March 2015, so the data has been labelled incorrectly. We shall therefore change all the values `MAR2015` to `MAR2016`.

```
data.loc[data["time"] == "MAR2015", "time"] = "MAR2016"
```

We can now change `time` into a categorical variable.

```
data["time"] = data["time"].astype("category")
```

**1** Read ‘S1 Text’ (<https://tinyco.re/9522240>), from the journal paper’s supporting information, which explains how the Store Price Survey data was collected.

- (a) In your own words, explain how the product information was recorded, and the measures that researchers took to ensure that the data was accurate and representative of the treatment group. What were some of the data collection issues that they encountered?
- (b) Instead of using the name of the store, each store was given a unique ID number (recorded as `store_id` on the spreadsheet). Verify that the number of stores in the dataset is the same as that stated in the ‘S1 Text’ (26). Similarly, each product was given a unique ID number (`product_id`). How many different products are in the dataset?

## PYTHON WALK-THROUGH 3.2

### *Counting the number of unique elements in a variable*

We use two methods here: the `unique` function to obtain a list of the unique elements of the variables of interest (`data["store_id"]` and `data["product_id"]`), then `len` to count how long the list is. We will create two variables, `no_stores` and `no_products`, that contain the number of stores and products respectively.

```
no_stores = len(data["store_id"].unique())
no_products = len(data["product_id"].unique())
print(f"The {no_stores} while the {no_products}.")
```

The no\_stores=26 while the no\_products=247.

Following the approach described in ‘S1 Text’, we will compare the variable price per ounce in US\$ cents (`price_per_oz_c`). We will look at what happened to prices in the two treatment groups before the tax (`time = DEC2014`) and after the tax (`time = JUN2015`):

- *treatment group one:* large supermarkets (`store_type = 1`)
- *treatment group two:* pharmacies (`store_type = 3`).

Before doing this analysis, we will use summary measures to see how many observations are in the treatment and control group, and how the two groups differ across some variables of interest. For example, if there are very few observations in a group, we might be concerned about the precision of our estimates and will need to interpret our results in light of this fact.

We will now create frequency tables containing the summary measures that we are interested in.

### 2 Create the following tables:

- (a) A frequency table showing the number (count) of store observations (store type) in December 2014 and June 2015, with `store_type` as the row variable and `time_period` as the column variable. For each store type, is the number of observations similar in each time period?
- (b) A frequency table showing the number of taxed and non-taxed beverages in December 2014 and June 2015, with `store_type` as the row variable and `taxed` as the column variable. (`Taxed` equals 1 if the sugar tax applied to that product, and 0 if the tax did not apply). For each store type, is the number of taxed and non-taxed beverages similar?
- (c) A frequency table showing the number of each product type (type), with `product_type` as the row variable and `time_period` as the column variable. Which product types have the highest number of

observations and which have the lowest number of observations?  
Why might some products have more observations than others?

### PYTHON WALK-THROUGH 3.3

#### *Creating frequency tables*

*Frequency table for store type and time period*  
We use the `pd.crosstab` function to perform a cross-tabulation count of `store_type` by `time`.

```
pd.crosstab(index=data["store_type"],
            columns=data["time"])
```

time	DEC2014	JUN2015	MAR2016
store_type			
Large supermarket	177	209	158
Small supermarket	407	391	327
Pharmacy	87	102	73
Gas station	73	96	75

There are fewer observations taken from gas stations and pharmacies and more from small supermarkets.

#### *Frequency table for store type and taxed*

Now we repeat the steps above to make a frequency table with `store_type` as the row variable and `taxed` as the column variable. We pass a list of dataframe columns to the `pd.crosstab` function because we also want separate values for each time period. To get the total row each row and column, we add `margins=True`.

```
pd.crosstab(
    index=data["store_type"],
    columns=[data["time"], data["taxed"]],
    margins=True
)
```

time	DEC2014		JUN2015		MAR2016		All
taxed	not taxed	taxed	not taxed	taxed	not taxed	taxed	not taxed
store_type							
Large supermarket	92	85	111	98	88	70	544
Small supermarket	196	211	192	199	154	173	1125
Pharmacy	44	43	52	50	36	37	262
Gas station	34	39	44	52	31	44	244
All	366	378	399	399	309	324	2175

#### *Frequency table for product type and time period*

Now we make a frequency table with product type (`type`) as the row variable and time period (`time`) as the column variable.

```
pd.crosstab(index=data["type"],
            columns=data["time"], margins=True)
```

time	DEC2014	JUN2015	MAR2016	All
type				
ENERGY	56	58	49	163
ENERGY-DIET	49	54	35	138
JUICE	70	64	52	186
JUICE DRINK	19	17	6	42
MILK	63	61	53	177
SODA	239	262	215	716
SODA-DIET	128	174	127	429
SPORT	11	16	12	39
SPORT-DIET	2	2	0	4
TEA	52	45	41	138
TEA-DIET	6	6	8	20
WATER	48	38	34	120
WATER-SWEET	1	1	1	3
All	744	798	633	2175

This table shows that there were no observations for the category **SPORT-DIET** in March 2016. As this is a drink which even in the other months has very few observations, it may be a product that is offered only in

one shop, and it is possible that this shop was not visited in March 2016. The product may also be a seasonal product that is not available in March. It is also likely that **WATER-SWEET** is offered in only one shop.

Besides counting the number of observations in a particular group, we can also calculate the mean by only using observations that satisfy certain conditions (known as the **conditional mean**). In this case, we are interested in comparing the mean price of taxed and untaxed beverages, before and after the tax.

### 3 Calculate and compare conditional means:

- (a) Create a table similar to Figure 3.1, showing the average price per ounce (in cents) for taxed and untaxed beverages separately, with **store\_type** as the row variable, and **taxed** and **time** as the column variables. To follow the methodology used in the journal paper, make sure to only include products that are present in all time periods, and non-supplementary products (**supp = 0**).
- (b) Without doing any calculations, summarize any differences or general patterns between December 2014 and June 2015 that you find in the table.
- (c) Would we be able to assess the effect of sugar taxes on product prices by comparing the average price of untaxed goods with that of taxed goods in any given period? Why or why not?

**conditional mean** An average of a variable, taken over a subgroup of observations that satisfy certain conditions, rather than all observations.

	Non-taxed	Taxed		
Store type	Dec 2014	Jun 2015	Dec 2014	Jun 2015
1				
3				

**Figure 3.1** The average price of taxed and non-taxed beverages, according to time period and store type.

## PYTHON WALK-THROUGH 3.4

### Calculating conditional means

Calculating conditional, or group, means is a common data cleaning operation you will encounter. It can be a bit confusing, but, in the typical case, we might just be asking what the mean is for each period (instead of the mean for all periods). Also, we must choose whether to put the mean back into our data with the *original index* (and so the original number of rows) or with a new index that reflects what the mean is conditional on. We'll now see one way to do this.

In order to identify products (`product_id`) that have observations for all three periods (`DEC2014`, `JUN2015` and `MAR2016`), we will first create a new variable called `period_test`, which takes the value `True` if we have observations in all periods for a product in a particular store, and `False` otherwise. These true/false variables are called Boolean variables.

This is going to be a complex operation with several steps. First, we'll use a `groupby` operation to group everything by `product_id` and `store_id`. Then we'll select the column we want (`time`). Then we'll use `transform`: this command returns a series with the same number of rows as the original index, so will fit nicely into our existing dataframe (rather than having the index set as the number of unique groups from the `groupby` operation). Within our `transform`, we'll ask for the number of unique (`nunique`) categories of time. For entries that exist, this value will be 1, 2, or 3 depending on whether the product is available in 1, 2, or 3 of the periods. For entries that don't exist, this value will be `NaN`, which means 'not a number'. So we then use `.fillna(0)` to set these to zero. Finally we compare this object against the known unique number of different periods (3) to get our `True` or `False` value; in this case, `True` if there are 3 periods, and `False` otherwise.

```
data.groupby(["product_id",
"store_id"])["time"].transform(
```

```
lambda x: x.nunique()
).fillna(0) == data["time"].nunique()

0      True
1      False
2      False
3      True
4      False
...
2170    True
2171    True
2172    True
2173    True
2174    True
Name: time, Length: 2175, dtype: bool
```

*Tip:* The `pandas` documentation provides many detailed examples of its extensive functionality, including the `transform` function.

Let's store this in the dataframe as a new column called `period_test`:

```
data["period_test"] = (
    data.groupby(["product_id",
"store_id"])["time"]
    .transform(lambda x: x.nunique())
    .fillna(0)
    == data["time"].nunique()
)
```

Now we can check this code worked by picking out the first unique `product_id`, `store_id` combination:

```
data.loc[(data["product_id"] == 29) &
         (data["store_id"] == 16), :]
```

	store_id	type	store_type	type2	size	price	price_per_oz	price_per_oz_c	taxed	supp	time	product_id
0	16	WATER	Small supermarket	NaN	33.8	1.69	0.050000	5.000000	not taxed	standard	DEC2014	29
744	16	WATER	Small supermarket	NaN	33.8	1.79	0.052959	5.295858	not taxed	standard	JUN2015	29
1542	16	WATER	Small supermarket	NaN	33.8	1.69	0.050000	5.000000	not taxed	standard	MAR2016	29

The table this produces is too wide for this page. It is available in full online (<https://tinyco.re/2230008>).

`period_test` reports `True` for this combination, and we can see that all three periods are present. Let's try another `product_id`, `store_id` combination:

```
data.loc[(data["product_id"] == 29) &
         (data["store_id"] == 10), :]
```

	store_id	type	store_type	type2	size	price	price_per_oz	price_per_oz_c	taxed	supp	time	product_id
69	10	WATER	Small supermarket	NaN	33.8	1.75	0.051775	5.177515	not taxed	standard	DEC2014	29
817	10	WATER	Small supermarket	NaN	33.8	1.99	0.058876	5.887574	not taxed	standard	JUN2015	29

The table this produces is too wide for this page. It is available in full online (<https://tinyco.re/2230008>).

In this case, the product-store cell is only there for two out of the three periods and so gets a `False` entry for its `period_test` column.

Now we can use the `period_test` variable to remove all products that have not been observed in all three periods. We'll also filter by `standard` in the

"`supp`" column. We will store the remaining products in a new dataframe, `data_c`.

```
data_c = data.loc[data["period_test"] &
                  (data["supp"] == "standard")]
data_c.head()
```

	store_id	type	store_type	type2	size	price	price_per_oz	price_per_oz_c	taxed	supp	time	product_id
0	16	WATER	Small supermarket	NaN	33.8	1.69	0.050000	5.000000	not taxed	standard	DEC2014	29
3	16	WATER	Small supermarket	NaN	33.8	1.69	0.050000	5.000000	not taxed	standard	DEC2014	38
5	16	MILK	Small supermarket	LOW FAT MILK	64.0	2.79	0.043594	4.359375	not taxed	standard	DEC2014	41
7	16	MILK	Small supermarket	WHOLE MILK	64.0	2.79	0.043594	4.359375	not taxed	standard	DEC2014	43
10	16	SODA	Small supermarket	NaN	20.0	1.89	0.094500	9.450000	taxed	standard	DEC2014	59

The table this produces is too wide for this page. It is available in full online (<https://tinyco.re/2230008>).

Now we can calculate the means of `price_per_oz` by grouping the data according to `store_type`, `taxed`, and `time`.

You may have noticed that we're stringing together multiple methods, with one called after another like `.method1(options).method2(options)`. This is a technique called *method chaining*: it just means chaining many successive methods together one after another.

Method chaining is a useful technique for data analysis, but it has advantages and disadvantages. If you know what the code is doing, and that it works already, it can be substantially easier to read than successive lines of logic using assignment via `=`. However, if something is going wrong within a long method chain, it can be hard to know where and why.

You can find a deeper dive into method chaining in the relevant section of Coding for Economists.

In the code below, we're going to use a lot of method chaining to get the data we want. The first operation in the chain groups all of the variables together by `taxed`, `store_type`, and `time`. The second operation produces some new columns via `.agg`. The column names, `avg_price` and `n`, are specified via the `("input`

`column", "function to be applied")` syntax. The new dataframe has the `groupby` variables as its index, so we reset the index to row numbers (`.reset_index`). We can then perform a `pivot` to transform the data into the shape that we want.

One point to note is that the `new_column_name = ("input column", "function to be applied")` syntax is special to `pandas`, so you likely won't see it in other packages. However, it is quite convenient in this case where we have quite a few bespoke operations and new column names being introduced all at once—and it's really helpful for method chaining!

```
table_res = (
    data_c.groupby(["taxed", "store_type",
    "time"])
    .agg(avg_price=("price_per_oz", "mean"),
    n=("price_per_oz", "count"))
    .reset_index()
    .pivot(index=["taxed", "store_type",
    "n"], columns="time", values="avg_price")
)
table_res
```

		time	DEC2014	JUN2015	MAR2016
taxed	store_type	n			
not taxed	<b>Large supermarket</b>	<b>36</b>	0.111920	0.114804	0.117015
	<b>Small supermarket</b>	<b>70</b>	0.136708	0.138165	0.133672
	<b>Pharmacy</b>	<b>18</b>	0.151958	0.160754	0.154353
	<b>Gas station</b>	<b>12</b>	0.169351	0.169641	0.170414
taxed	<b>Large supermarket</b>	<b>36</b>	0.156174	0.169297	0.166843
	<b>Small supermarket</b>	<b>101</b>	0.158510	0.159952	0.154925
	<b>Pharmacy</b>	<b>18</b>	0.181818	0.190788	0.186303
	<b>Gas station</b>	<b>22</b>	0.194117	0.203367	0.192376

Use the S3 Table in the journal paper to check how closely your summary data match those in the paper. You should find that your results for Large

Supermarkets and Pharmacies match, but the other store types have discrepancies. In Python walk-through 3.5, we will discuss these differences in more detail.

In order to make a before-and-after comparison, we will make a chart similar to Figure 2 (<https://tinyco.re/8127041>) in the journal paper, to show the change in prices for each store type.

#### 4 Using your table from Question 3:

- (a) Calculate the change in the mean price after the tax (price in June 2015 minus price in December 2014) for taxed and untaxed beverages, by store type.
- (b) Using the values you calculated in Question 4(a), plot a column chart to show this information (as done in Figure 2 of the journal paper) with store type on the horizontal axis and price change on the vertical axis. Label each axis and data series appropriately. You should get the same values as shown in Figure 2.

### PYTHON WALK-THROUGH 3.5

*Making a column chart to compare two groups*

*Calculate price differences by store type*

Let's calculate the two price differences (June 2015 minus December 2014 and March 2016 minus December 2014), and store them as `d1` and `d2`, respectively:

```
table_res["d1"] = table_res["JUN2015"] -
    table_res["DEC2014"]
table_res["d2"] = table_res["MAR2016"] -
    table_res["DEC2014"]
table_res
```

			time	DEC2014	JUN2015	MAR2016	d1	d2
taxed	store_type	n						
	Large supermarket	36	0.111920	0.114804	0.117015	0.002884	0.005096	
not taxed	Small supermarket	70	0.136708	0.138165	0.133672	0.001458	-0.003036	
	Pharmacy	18	0.151958	0.160754	0.154353	0.008796	0.002395	
	Gas station	12	0.169351	0.169641	0.170414	0.000290	0.001063	
	Large supermarket	36	0.156174	0.169297	0.166843	0.013122	0.010669	
taxed	Small supermarket	101	0.158510	0.159952	0.154925	0.001442	-0.003585	
	Pharmacy	18	0.181818	0.190788	0.186303	0.008970	0.004485	
	Gas station	22	0.194117	0.203367	0.192376	0.009251	-0.001740	

*Plot a column chart for average price changes*

To display `d1` and `d2` in a column chart, we will use the `geom_bar` function in the `lets_plot` package, which we loaded at the start of this project. We'll use `stat="identity"` to get `lets_plot` to use the numbers in the dataframe rather than the count of entries in the dataframe (which is its default). `position_dodge()` causes the two types we are looking at in `"taxed"` to appear side-by-side rather than overlaid.

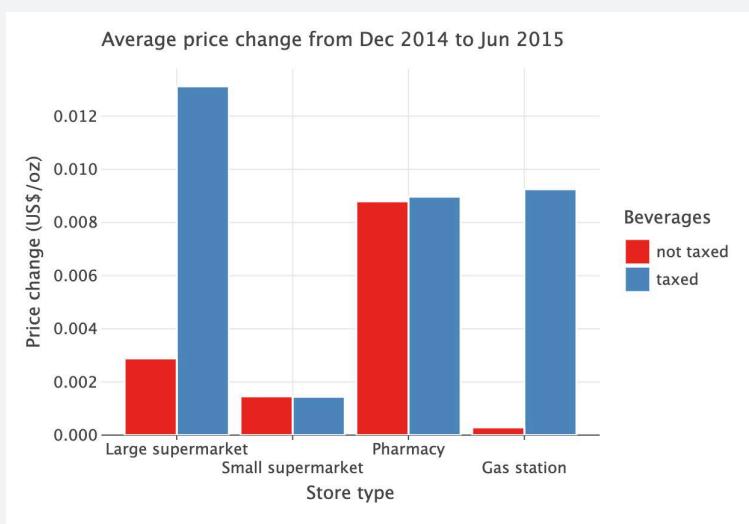
Let's start with displaying the average price change from December 2014 to June 2015 (which is stored in `d1`).

```
(  
  ggplot(table_res.reset_index(), aes(fill="taxed",  
y="d1", x="store_type"))  
  + geom_bar(stat="identity", position=position_dodge())
```

```

+ labs(
  y="Price change (US$/oz)",
  x="Store type",
  title="Average price change from Dec 2014 to Jun
2015",
  fill="Beverages",
)
)

```



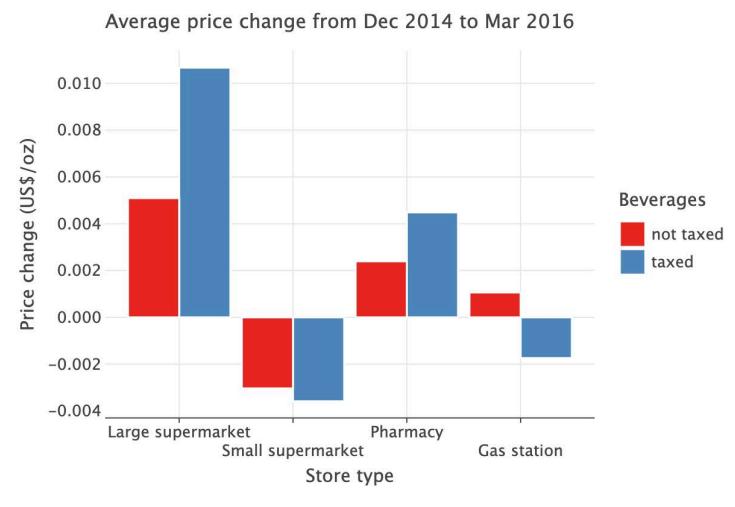
**Figure 3.2** Average price change from December 2014 to June 2015.

Now we do the same for the price change from Dec 2014 to Mar 2016:

```

(
  ggplot(table_res.reset_index(), aes(fill="taxed",
y="d2", x="store_type"))
  + geom_bar(stat="identity", position=position_dodge())
  + labs(
    y="Price change (US$/oz)",
    x="Store type",
    title="Average price change from Dec 2014 to Mar
2016",
    fill="Beverages",
  )
)

```



**Figure 3.3** Average price change from December 2014 to March 2016.

**statistically significant** When a relationship between two or more variables is unlikely to be due to chance, given the assumptions made about the variables (for example, having the same mean). Statistical significance does not tell us whether there is a causal link between the variables.

To assess whether the difference in mean prices before and after the tax could have happened by chance due to the samples chosen (and there are no differences in the population means), we could calculate the p-value. (Here, ‘population means’ refer to the mean prices before/after the tax that we would calculate if we had all prices for all stores in Berkeley.) The authors of the journal article calculate p-values, and use the idea of statistical significance to interpret them. Whenever they get a p-value of less than 5%, they conclude that the assumption of no differences in the population is unlikely to be true: they say that the price difference is **statistically significant**. If they get a p-value higher than 5%, they say that the difference is not statistically significant, meaning that they think it could be due to chance variation in prices.

Using a particular cutoff level for the p-value, and concluding that a result is only statistically significant if the p-value is below the cutoff, is common in statistical studies, and 5% is often used as the cutoff level. But this approach has been criticized recently by statisticians and social scientists. The main criticisms raised are that any cutoffs are arbitrary. Instead of using a cutoff, we prefer to calculate p-values and use them to assess the strength of the evidence against our assumption that there are no differences in the population means. Whether the statistical evidence is strong enough for us to draw a conclusion about a policy, such as a sugar tax, will always be a matter of judgement.

According to the journal paper, the p-value is 0.02 for large supermarkets, and 0.99 for pharmacies.

- 5 Based on these p-values and your chart from Question 4, what can you conclude about the difference in means? (*Hint:* You may find the discussion in Part 2.3 (page 174) helpful.)

## EXTENSION PYTHON WALK-THROUGH 3.6

### *Calculating the p-value for price changes*

In this walk-through, we show the calculations required to obtain the p-values in Table S3 of the Silver et al. (2017) paper. Since the p-values are already provided, this walk-through is only for those who want to see how these p-values were calculated. For the categories of Large supermarkets and Pharmacies, we conduct a hypothesis test, which tests the null hypothesis that the price difference between June 2015 and December 2014 (and March 2016 and December 2014) for the taxed and untaxed beverages in the two store types could be due to chance.

We are interested in whether the difference in average price between `JUN2015` and `DEC2014` (or `MAR2016` and `DEC2014`) for one group (say, Large supermarket and taxed beverages) is zero (i.e. there is no difference in the means of the two populations). Note that we are dealing with paired observations (the same product in both time periods).

Let's use the price difference between June 2015 and December 2014 in Large supermarkets for taxed beverages as an example. First, we extract the prices for both periods (the vectors `p_1` and `p_2`) and then calculate the difference, element by element (stored as `d_t`). Because *all operations assume the same index*, we must also reset the indexes of these sets of values that have been drawn from different rows (dropping the original indexes as we go).

In the code below, we're going to use a different way to cut a dataframe: `query`. Previously, to filter down a dataframe by the values in its columns, we've used `.loc` with the filters we wanted. For example, to filter rows with `value_a` in `column1` and `value_b` in `column2`:

```
df.loc[(df[column1] == "value_a") &
       (df[column2]== "value_b"), :]
```

This syntax can result in very long lines of code. `query` allows you to write something that is shorter and is composed of a single string (which sometimes comes in handy). The syntax for `query` uses `col_name == 'value_name'` with `&` to join multiple filters together. Note that if the value used to filter the column is numeric instead of string or category, you do not need the inner quotes, for example, `col_name = 85`.

```
p_1 = data_c.query(
    "store_type == 'Large supermarket' & taxed == 'taxed'"
    & time == 'DEC2014'"
)[["price_per_oz"]].reset_index(drop=True)
p_2 = data_c.query(
```

```

    "store_type == 'Large supermarket' & taxed == 'taxed'
    & time == 'JUN2015'"
)[["price_per_oz"]].reset_index(drop=True)

# Price difference for taxed products in large supermarkets
d_t = p_2 - p_1
d_t.head()

```

```

0    0.010000
1    0.010059
2    0.009600
3    0.010000
4    0.010059
Name: price_per_oz, dtype: float64

```

All three new variables are vectors with 36 elements. For `d_t` to correctly represent the price difference for a particular product in a particular store, we need to be certain that each element in both vectors corresponds to the same product in the same store.

To check that the elements match, we will extract the store and product IDs along with the prices, and save it as `alt`. We will then compare the ordering in the `DEC2014` column (`p1_alt`) and the `JUN2015` column (`p2_alt`).

```

alt = data_c.loc[
    (data_c["store_type"] == "Large supermarket") &
    (data_c["taxed"] == "taxed"), :
].pivot(index=["store_id", "product_id"], columns="time",
values="price_per_oz")
alt.head()

```

		time	DEC2014	JUN2015	MAR2016
store_id	product_id				
25	59		0.099500	0.109500	0.109500
	60		0.032396	0.042456	0.042456
	153		0.135200	0.144800	0.144800
	189		0.094500	0.104500	0.109500
	190		0.029438	0.039497	0.039497

Now we can simply subtract column `p_2_alt` from `p_1_alt`, preserving the index. We'll save this difference as `d_t_alt`.

```
p_1_alt, p_2_alt = alt["DEC2014"], alt["JUN2015"]
d_t_alt = p_2_alt - p_1_alt
d_t_alt.head()
```

```
store_id  product_id
25        59          0.010000
           60          0.010059
           153         0.009600
           189         0.010000
           190         0.010059
dtype: float64
```

Note that in this case, we can also see which combination produced this particular value of the `price_per_oz`. This approach guarantees identical ordering because the subtraction would fail if the indices were different.

The average value of the price difference is 0.1312, and our task is to evaluate whether this is likely to be due to sampling variation (given the assumption that there is no difference between the populations) or not. To do this, we can use the `ttest` function, which provides the associated p-value.

```
# Recognise that the differences come from paired samples.
pg.ttest(p_2, p_1, paired=True)
```

	T	dof	alternative	p-val	CI95%	cohen-d	BF10	power
<b>ttest</b>	4.768132	35	two-sided	0.000032	[0.01, 0.02]	0.079958	699.454	0.075305

Alternatively, we can calculate the respective test statistic manually, which also requires us to calculate the standard error of this value:

```
d_t_alt.mean() / np.sqrt((d_t_alt.var() / len(d_t_alt)))
```

**4.768131783937789**

Yet another alternative is to use the `ttest` function directly on `d_t`.

To compare this result to the journal paper, look at the extract from Table S3 (the section on Large supermarkets) shown in Figure 3.4 below. The cell with **\*\*** shows the mean price difference of 1.31 cents (\$0.0131).

Large supermarkets (n = 6)	Taxed beverage price (36 sets)	Untaxed beverage price (36 sets)	Taxed – untaxed difference
	cents/oz	cents/oz	cents/oz
Round 1: December 2014	15.62	11.19	
Round 2: June 2015	16.93	11.48	
Round 3: March 2016	16.68	11.70	
Mean change (March 2016–Dec 2014)	1.07, (p=0.01)	0.51, (p=0.01)	0.56, (p=0.22)
Mean change (June 2015–Dec 2014)	1.31, (p<0.001)**	0.29, (p=0.08)	1.02, (p=0.002)

**Figure 3.4** Table S3 in Silver et al. (2017), showing means and confidence intervals. n = number of stores of each type.

In our test output, we get a very small p-value (0.000032) which in the table is indicated by the double asterisk.

Tests for other store types are calculated similarly, by changing the data extracted to `p_1` and `p_2`. Let's do that for one more example: the price difference between June 2015 and December 2014 in Large supermarkets for untaxed beverages.

```
alt_nt = data_c.loc[
    (data_c["store_type"] == "Large supermarket") &
    (data_c["taxed"] == "not taxed"), :
].pivot(index=["store_id", "product_id"], columns="time",
values="price_per_oz")
d_nt_alt = alt_nt["JUN2015"] - alt_nt["DEC2014"]
print(f"Mean difference is {100*d_nt_alt.mean():.2f} cents/oz")
pg.ttest(d_nt_alt, y=0)
```

Mean difference is **0.29** cents/oz

	T	dof	alternative	p-val	CI95%	cohen-d	BF10	power
<b>ttest</b>	1.817855	35	two-sided	0.077655	[−0.0, 0.302976	0.788	0.423962	0.01]

You should be able to recognise the mean difference and the p-value in the excerpt of Table S3 provided in Figure 3.4.

Let's also replicate the last section of Table S3, which shows the difference between the price changes in taxed and untaxed products, that is, we want to know whether `d_t_alt` and `d_nt_alt` have different means. We will apply the two sample hypothesis tests, but this time for unpaired data, as the products differ across samples.

```
pg.ttest(d_t_alt, d_nt_alt)
```

	T	dof	alternative	p-val	CI95%	cohen-d	BF10	power
<b>ttest</b>	3.222715	70	two-sided	0.001929	[0.0, 0.02]	0.759601	17.601	0.888432

Again you should be able to identify the corresponding entries in Table S3 shown in Figure 3.4. The main entry in the table is 1.02, indicating that the means of the two groups differ by 1.02 cents. This is confirmed in our calculations, as \$0.01312 – \$0.00288 is about \$0.0102 or 1.02 cents. The p-value of 0.002 is also the same as the one in Table S3.

## PART 3.2 BEFORE-AND-AFTER COMPARISONS WITH PRICES IN OTHER AREAS

When looking for any price patterns, it is possible that the observed changes in Berkeley were not solely due to the tax, but instead were also influenced by other events that happened in Berkeley and in neighbouring areas. To investigate whether this is the case, the researchers conducted another differences-in-differences analysis, using a different treatment and control group:

- *The treatment group:* Beverages in Berkeley
- *The control group:* Beverages in surrounding areas.

The researchers collected price data from stores in the surrounding areas and compared them with prices in Berkeley. If prices changed in a similar way in nearby areas (which were not subject to the tax), then what we observed in Berkeley may not be primarily a result of the tax. We will be using the data the researchers collected to make our own comparisons.

Download the following files:

- The Berkeley Point-of-Sale Stata file on the Global Food Research Program's website (<https://tinyco.re/7121674>), containing the price data they collected, including information on the date (year and month), location (Berkeley or non-Berkeley), beverage group (soda, fruit drinks, milk substitutes, milk and water), and the average price for that month. Stata is another popular statistical software package, and the data is provided as a `.dta` file.
- 'S5 Table' (<https://tinyco.re/7724734>) which compares the neighbourhood characteristics of the Berkeley and non-Berkeley stores.

- 1** Based on ‘S5 Table’, do you think the researchers chose suitable comparison stores? Why or why not?

We will now plot a line chart similar to Figure 3 in the journal paper, to compare prices of similar goods in different locations and see how they have changed over time. To do this, we will need to summarize the data so that there is one value (the mean price) for each location and type of good in each month.

- 2** Assess the effects of a tax on prices:

- (a) Create a table similar to the one provided in Figure 3.5 to show the average price in each month for taxed and non-taxed beverages, according to location. Use ‘year and month’ as the row variables, and ‘tax’ and ‘location’ as the column variables. (*Hint:* You may find Python walk-through 3.4 (page 300) helpful.)
- (b) Plot the four columns of your table on the same line chart, with average price on the vertical axis and time (months) on the horizontal axis. Describe any differences you see between the prices of non-taxed goods in Berkeley and those outside Berkeley, both before the tax (January 2013 to December 2014) and after the tax (March 2015 onwards). Do the same for prices of taxed goods.
- (c) Based on your chart, is it reasonable to conclude that the sugar tax had an effect on prices?

Year/Month	Non-taxed		Taxed	
	Berkeley	Non-Berkeley	Berkeley	Non-Berkeley
January 2013				
February 2013				
March 2013				
...				
December 2013				
January 2014				
...				
February 2016				

**Figure 3.5** The average price of taxed and non-taxed beverages, according to location and month.

## PYTHON WALK-THROUGH 3.7

### Importing data from a Stata file and plotting a line chart

Import data and create a table of average monthly prices

To import data from a Stata file (.dta format) we need the pandas package's pd.read\_stata function.

Remember to download the data from this link and save the file with extension .dta in the data/ folder of your working directory.

```
path_to_stata_data = Path("data/public_use_weighted_prices2.dta")
posd = pd.read_stata(path_to_stata_data)
posd.head()
```

	year	quarter	month	location	beverage_group	tax	price	under_report
0	2013.0	1.0	1.0	Berkeley	soda	Non-taxed	4.853399	NaN
1	2013.0	1.0	1.0	Non-Berkeley	soda	Non-taxed	3.510271	NaN
2	2013.0	1.0	1.0	Non-Berkeley	soda	Non-taxed	3.889572	NaN
3	2013.0	1.0	1.0	Berkeley	soda	Taxed	3.682805	NaN
4	2013.0	1.0	1.0	Non-Berkeley	soda	Taxed	3.516437	NaN

For each month and location (Berkeley or Non-Berkeley), there are prices for a variety of beverage categories, and we know whether the category is taxed or not. For any particular time-location-tax status combination, we want the average price of all products.

To make the summary table, we use method chaining again:

```
table_test = (
    posd.groupby(["year", "month",
    "location", "tax"])["price"]
    .agg("mean")
    .reset_index()
)
table_test
```

	year	month	location	tax	price
0	2013.0	1.0	Berkeley	Non-taxed	5.722480
1	2013.0	1.0	Berkeley	Taxed	8.692803
2	2013.0	1.0	Non-Berkeley	Non-taxed	5.348640
3	2013.0	1.0	Non-Berkeley	Taxed	7.991574
4	2013.0	2.0	Berkeley	Non-taxed	5.806468
...	...	...	...	...	...
151	2016.0	2.0	Non-Berkeley	Taxed	8.729508
152	2016.0	3.0	Berkeley	Non-taxed	NaN
153	2016.0	3.0	Berkeley	Taxed	NaN
154	2016.0	3.0	Non-Berkeley	Non-taxed	NaN
155	2016.0	3.0	Non-Berkeley	Taxed	NaN

We'll now prepare the data for plotting in a line chart.

First, we reset the index so that it is just composed of row numbers. We'd also like to add a column for time, recorded as a time series (in this case, a sequence of years and months).

```
table_test["date"] = (
    table_test.apply(
        lambda row: pd.to_datetime(
            str(int(row["year"])) + "-" +
            str(int(row["month"]))), format="%Y-%m"
```

```

),
axis=1,
)
+ pd.offsets.MonthEnd()
)
table_test.head()

```

	year	month	location	tax	price	date
0	2013.0	1.0	Berkeley	Non-taxed	5.722480	2013-01-31
1	2013.0	1.0	Berkeley	Taxed	8.692803	2013-01-31
2	2013.0	1.0	Non-Berkeley	Non-taxed	5.348640	2013-01-31
3	2013.0	1.0	Non-Berkeley	Taxed	7.991574	2013-01-31
4	2013.0	2.0	Berkeley	Non-taxed	5.806468	2013-02-28

Here, we create a variable called `date` by passing an integer version of the year and month together as a string to `pd.to_datetime`, specifying the date format as YYYY-MM format (`%Y-%m`). This code produces a date at the start of each month; `pd.offsets.MonthEnd()` finds the last day in the month.

### Plot a line chart

Now we can create a line chart.

```

date_one = pd.to_datetime("2015-01-01")
date_two = pd.to_datetime("2015-03-01")

# Rename 'date' as 'Date' for the horizontal
# axis label
table_test =
table_test.rename(columns={'date': 'Date'})
(
    ggplot(table_test, aes(x="Date",
y="price", color="tax", linetype="location"))
    + geom_vline(xintercept=date_one,
color="gray")
    + geom_vline(xintercept=date_two,
color="gray")
    + geom_line(size=1)
    + geom_text(
        x=date_one -
pd.DateOffset(months=3), y=5,
label="Pre-tax", color="black"
    )
    + geom_text(
        x=date_two +

```

```

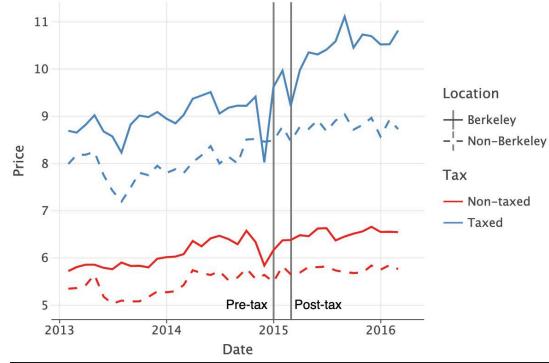
+ pd.offsets.MonthEnd()
)
table_test.head()

```

In the code above, we used `apply` and `lambda` again, both methods that we've seen before. Remember that, given an axis (`axis=`), the `apply` method applies a function to every element. Here we apply it to every row (`axis=1`) because the date column we want to create relies on combining information from different columns for each row.

```

pd.DateOffset(months=3), y=5,
label="Post-tax", color="black"
)
+ labs(y="Price", color="Tax",
linetype="Location")
)
```



**Figure 3.6** Average price of taxed and non-taxed beverages in Berkeley and non-Berkeley areas.

How strong is the evidence that the sugar tax affected prices? According to the journal paper, when comparing the mean Berkeley and non-Berkeley price of sugary beverages after the tax, the p-value is smaller than 0.00001, and it is 0.63 for non-sugary beverages after the tax.

- 3** What do the p-values tell us about the difference in means and the effect of the sugar tax on the price of sugary beverages? (*Hint:* You may find the discussion in Part 2.3 (page 174) helpful.)

The aim of the sugar tax was to decrease consumption of sugary beverages. Figure 3.7 shows the mean number of calories consumed and the mean volume consumed before and after the tax. The researchers reported the p-values for the difference in means before and after the tax in the last column.

- 4** Based on Figure 3.7, what can you say about consumption behaviour in Berkeley after the tax? Suggest some explanations for the evidence.
- 5** Read the ‘Limitations’ in the ‘Discussions’ section (<https://tinyco.re/6616217>) of the paper and discuss the strengths and limitations of this study. How could future studies on the sugar tax in Berkeley address these problems? (Some issues you may want to discuss are: the number of stores observed, number of people surveyed, and the reliability of the price data collected.)

Usual intake	Pre-tax (Nov-Dec 2014), n = 623	Post-tax (Nov-Dec 2015), n = 613	Pre-tax-post-tax difference
<b>Caloric intake (kilocalories/capita/day)</b>			
Taxed beverages	45.1	38.7	-6.4, p = 0.56
Non-taxed beverages	115.7	147.6	31.9, p = 0.04
<b>Volume of intake (grams/capita/day)</b>			
Taxed beverages	121.0	97.0	-24.0, p = 0.24
Non-taxed beverages	1,839.4	1,896.5	57.1, p = 0.22

Models account for age, gender, race/ethnicity, income level, and educational attainment. n is the sample size at each round of the survey after excluding participants with missing values on self-reported race/ethnicity, age, education, income, or monthly intake of sugar-sweetened beverages.

Lynn D. Silver, Shu Wen Ng, Suzanne Ryan-Ibarra, Lindsey Smith Taillie, Marta Induni, Donna R. Miles, Jennifer M. Poti, and Barry M. Popkin. 2017. Table 1 in ‘Changes in prices, sales, consumer spending, and beverage consumption one year after a tax on sugar-sweetened beverages in Berkeley, California, US: A before-and-after study’ (<https://tinyco.re/8147535>). PLoS Med 14 (4): e1002283.

**Figure 3.7** Changes in prices, sales, consumer spending, and beverage consumption one year after a tax on sugar-sweetened beverages in Berkeley, California, US: A before-and-after study.

- 6** Suppose that you have the authority to conduct your own sugar tax natural experiment in two neighbouring towns, Town A and Town B. Outline how you would conduct the experiment to ensure that any changes in outcomes (prices, consumption of sugary beverages) are due to the tax and not due to other factors. (*Hint:* think about what factors you need to hold constant.)